

JIEUN HWANG

3033297165

```
In [150]: # Python 2 and 3 support
from __future__ import division, print_function, unicode_literals

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline

# Hide warnings
import warnings
warnings.filterwarnings('ignore')

import tensorflow as tf
```

```
In [151]: from datetime import datetime
import os
import pathlib

t = datetime.utcnow().strftime("%Y%m%d%H%M%S")
log_dir = "tf_logs"
logd = "/tmp/{}/r{}/".format(log_dir, t)

# Then every time you have specified a graph run:
# file_writer = tf.summary.FileWriter(logdir, tf.get_default_graph())

# Make directory if it doesn't exist

from pathlib import Path
home = str(Path.home())

logdir = os.path.join(os.sep, home, logd)

if not os.path.exists(logdir):
    os.makedirs(logdir)
```

```
In [152]: # TensorBoard Graph visualizer in notebook
import numpy as np
from IPython.display import clear_output, Image, display, HTML

def strip_consts(graph_def, max_const_size=32):
    """Strip large constant values from graph_def."""
    strip_def = tf.GraphDef()
    for n0 in graph_def.node:
        n = strip_def.node.add()
        n.MergeFrom(n0)
        if n.op == 'Const':
            tensor = n.attr['value'].tensor
            size = len(tensor.tensor_content)
            if size > max_const_size:
                tensor.tensor_content = "<stripped %d bytes>"%size
    return strip_def

def show_graph(graph_def, max_const_size=32):
    """Visualize TensorFlow graph."""
    if hasattr(graph_def, 'as_graph_def'):
        graph_def = graph_def.as_graph_def()
    strip_def = strip_consts(graph_def, max_const_size=max_const_size)
    code = """
        <script src="//cdnjs.cloudflare.com/ajax/libs/polymer/0.3.3/platform.js"></script>
        <script>
            function load() {{
                document.getElementById("{id}").pbtxt = {data};
            }}
        </script>
        <link rel="import" href="https://tensorboard.appspot.com/tf-graph-basic.build.html" onload=load(
    )>

        <div style="height:600px">
            <tf-graph-basic id="{id}"></tf-graph-basic>
        </div>
    """.format(data=repr(str(strip_def)), id='graph'+str(np.random.rand()))

    iframe = """
        <iframe seamless style="width:1200px;height:620px;border:0" srcdoc="{></iframe>
    """.format(code.replace("'", '&quot;'))
    display(HTML(iframe))
```

```
In [153]: # Read in input data
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
# contains info
import tensorflow.examples.tutorials.mnist.mnist as mnist_info
```

```
Extracting MNIST_data/train-images-idx3-ubyte.gz
Extracting MNIST_data/train-labels-idx1-ubyte.gz
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
```

```
In [154]: # Read input_data (not as one_hot)
from tensorflow.examples.tutorials.mnist import input_data

# new folder
mnist = input_data.read_data_sets("/tmp/data/")

# Assign them to values
X_train = mnist.train.images
X_test = mnist.test.images
y_train = mnist.train.labels.astype("int")
y_test = mnist.test.labels.astype("int")
```

```
Extracting /tmp/data/train-images-idx3-ubyte.gz
Extracting /tmp/data/train-labels-idx1-ubyte.gz
Extracting /tmp/data/t10k-images-idx3-ubyte.gz
Extracting /tmp/data/t10k-labels-idx1-ubyte.gz
```

```
In [155]: # Define hyperparameters and input size 784
n_inputs = 28*28 # MNIST

# Three hidden layers of size 300,200,100
n_hidden1 = 300
n_hidden2 = 200
n_hidden3 = 100

# Output layer of size 10 ( to classify digits 0-9)
n_outputs = 10
```

```
In [156]: # Reset graph
tf.reset_default_graph()
```

```
In [157]: # Placeholders for data (inputs and targets)
X = tf.placeholder(tf.float32, shape=(None, n_inputs), name="X")
y = tf.placeholder(tf.int64, shape=(None), name="y")
```

```
In [158]: # Define neuron layers (ReLU in hidden layers)
# We'll take care of Softmax for output with loss function

def neuron_layer(X, n_neurons, name, activation=None):
    # X input to neuron
    # number of neurons for the layer
    # name of layer
    # pass in eventual activation function

    with tf.name_scope(name):
        n_inputs = int(X.get_shape()[1])

        # initialize weights to prevent vanishing / exploding gradients
        stddev = 2 / np.sqrt(n_inputs)
        init = tf.truncated_normal((n_inputs, n_neurons), stddev=stddev)

        # Initialize weights for the layer
        W = tf.Variable(init, name="weights")
        # biases
        b = tf.Variable(tf.zeros([n_neurons]), name="bias")

        # Output from every neuron
        Z = tf.matmul(X, W) + b
        if activation is not None:
            return activation(Z)
        else:
            return Z
```

```
In [159]: # Define the hidden layers and Use a dropout ratio of 10% on all hidden layers
with tf.name_scope("dnn"):
    hidden1 = neuron_layer(X, n_hidden1, name="hidden1",
                           activation=tf.nn.relu)
    hidden1 = tf.layers.dropout(hidden1, rate=0.1)

    hidden2 = neuron_layer(hidden1, n_hidden2, name="hidden2",
                           activation=tf.nn.relu)
    hidden2 = tf.layers.dropout(hidden2, rate=0.1)
    hidden3 = neuron_layer(hidden2, n_hidden3, name="hidden3",
                           activation=tf.nn.relu)
    hidden3 = tf.layers.dropout(hidden3, rate=0.1)
    logits = neuron_layer(hidden3, n_outputs, name="outputs")
```

```
In [160]: # Define loss function (that also optimizes Softmax for output) and Use cross entropy loss

with tf.name_scope("loss"):
    # logits are from the last output of the dnn
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y,
                                                              logits=logits)

    loss = tf.reduce_mean(xentropy, name="loss")
```

```
In [161]: # Training step with Gradient Descent

learning_rate = 0.01

with tf.name_scope("train"):
    optimizer = tf.train.GradientDescentOptimizer(learning_rate)
    training_op = optimizer.minimize(loss)
```

```
In [162]: # Evaluation to see accuracy
with tf.name_scope("eval"):
    correct = tf.nn.in_top_k(logits, y, 1)
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
```

```
In [163]: show_graph(tf.get_default_graph())
```

Fit to screen

Run

Upload

Choose File

Color

Structure

color: same substructure
gray: unique substructure

Graph

(* = expandable)

Namespace*

OpNode

Unconnected series*

Connected series*

Constant

Summary

Dataflow edge

Control dependency edge

Reference edge

Main Graph

The diagram is a computational graph with the following components:

- Nodes:**
 - train**: A rounded rectangle node at the top right.
 - loss**: A rounded rectangle node on the left.
 - eval**: A rounded rectangle node in the center.
 - dnn**: A rounded rectangle node below 'eval'.
 - y**: An oval node below 'loss'.
 - x**: An oval node at the bottom right.
- Edges:**
 - A curved dataflow edge from **train** to **loss**.
 - A curved dataflow edge from **train** to **eval**.
 - A curved dataflow edge from **train** to **x**.
 - A straight dataflow edge from **loss** to **y**.
 - A straight dataflow edge from **loss** to **dnn**.
 - A straight dataflow edge from **y** to **dnn**.
 - A curved dataflow edge from **eval** to **dnn**.
 - A curved dataflow edge from **dnn** to **x**.

http://localhost:8888/nbconvert/html/data-x/06b-tools-tensorflow/HW10.ipynb?download=false

Page 4 of 6

```
In [138]: init = tf.global_variables_initializer()
saver = tf.train.Saver()

n_epochs = 10
batch_size = 50

with tf.Session() as sess:
    init.run()
    for epoch in range(n_epochs):
        for iteration in range(mnist.train.num_examples // batch_size):
            X_batch, y_batch = mnist.train.next_batch(batch_size)
            sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
            acc_train = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
            acc_val = accuracy.eval(feed_dict={X: mnist.validation.images,
                                                y: mnist.validation.labels})
            print(epoch, "Train accuracy:", acc_train, "Val accuracy:", acc_val)

    save_path = saver.save(sess, "./my_model_final.ckpt") # save model

0 Train accuracy: 0.98 Val accuracy: 0.928
1 Train accuracy: 0.96 Val accuracy: 0.9434
2 Train accuracy: 0.94 Val accuracy: 0.9524
3 Train accuracy: 1.0 Val accuracy: 0.9574
4 Train accuracy: 0.98 Val accuracy: 0.9604
5 Train accuracy: 0.98 Val accuracy: 0.963
6 Train accuracy: 0.94 Val accuracy: 0.9666
7 Train accuracy: 1.0 Val accuracy: 0.9682
8 Train accuracy: 0.94 Val accuracy: 0.971
9 Train accuracy: 0.96 Val accuracy: 0.9714
```

```
In [139]: with tf.Session() as sess:
    saver.restore(sess, "./my_model_final.ckpt") # or better, use save_path
    X_new_scaled = mnist.test.images[:20]
    Z = logits.eval(feed_dict={X: X_new_scaled})
    y_pred = np.argmax(Z, axis=1)

print("Predicted classes:", y_pred)
print("Actual classes:   ", mnist.test.labels[:20])

INFO:tensorflow:Restoring parameters from ./my_model_final.ckpt
Predicted classes: [7 2 1 0 4 1 4 9 6 9 0 6 9 0 1 5 9 7 3 4]
Actual classes:    [7 2 1 0 4 1 4 9 5 9 0 6 9 0 1 5 9 7 3 4]
```

```
In [140]: show_graph(tf.get_default_graph())
```

Fit to screen

Run

Upload

Choose File

Color

Structure

color: same substructure

gray: unique substructure

Graph

(* = expandable)

Namespace*

OpNode

Unconnected series*

Connected series*

Constant


Summary

Dataflow edge

Control dependency edge

Reference edge

Main Graph



Auxiliary nod



Overall, I got Train accuracy: 0.96 Val accuracy: 0.9714

